

Linux Kernel Programming

Xiaoguang Wang

About me



Assistant Professor @ UIC CS

Email: xgwang9@uic.edu (preferred way to communicate)

Office: SEO 1331

Homepage: <https://xiaoguang.wang/>

My research interests

Systems software (operating systems, compiler, runtime, ...)

- How to efficiently execute programs on CPUs of heterogeneous architecture (e.g., SmartNIC, IoT/Edge)?

Software security

- How to leverage LLMs (e.g., ChatGPT) to better discover unknown bugs?
- Can we *automatically* detect and fix bugs?
- How to compartmentalize large-scale software to prevent bugs from spreading?

What is the Linux Kernel?

One of operating system kernels

- e.g., Windows, FreeBSD, OSX, etc.

What does an OS do for you?

- **Abstract** the hardware for convenience and portability
- **Multiplex** the hardware among multiple applications
- **Isolate** applications to contain bugs
- Allow **sharing** among applications


View: layered organization

User: applications (e.g., vim and gcc)

Kernel: file system, process, etc.

Hardware: CPU, mem, disk, etc.

→ Interface between layers



View: core services

- Process
- Memory
- File contents
- Directories and file names
- Security
- Many others: users, IPC, network, time, terminals, etc.

→ Abstraction for applications

Example: system calls

Interface: applications talk to an OS via system calls

Abstraction: process and file descriptor

```
fd = open("out", 1);  
write(fd, "hello\n", 6);  
pid = fork();
```

Why is Linux kernel interesting?

OS design deals with conflicting goals and trade-offs


- Efficient yet portable
- Powerful yet simple
- Isolated yet interactable
- General yet performant

Open problems: multi-core (scalability), performance, heterogeneous devices, security, ...

How does a state-of-the-art OS deal with above issues?

- Hack the Linux kernel!

Why is Linux kernel interesting?



Extremely large software project

- more than 30 ~~25~~ million lines of code
- 7,500 ~~4,600~~ lines of code are added every day!

<https://www.linuxfoundation.org/resources/publications/linux-kernel-history-report-2020>

Why is Linux kernel interesting?

Very fast development cycles

- release about every 70 days
- 13,000 patches / release

One of the most well-written/designed/maintained C code projects

More here:

- [Linux Foundation Annual Report 2021](#)

Linux is eating the World

71% of smartphones and tablets run Linux (Android), as of Dec/2023

- iOS: 28%

98% of top 1 million web servers run Linux

99% of super computers run Linux

SpaceX: [From Earth to orbit with Linux and SpaceX](#)

Ref: [Use share of operating systems](#)

It is good for your job search

Contributions from unpaid developers had been in slow decline

- 14.6% (2012) → 13.6% (2013) → 11.8% (2014) → 7.7% (2015)

Why?

- “There are many possible reasons for this decline, but, arguably, the most plausible of those is quite simple: **Kernel developers are in short supply, so anybody who demonstrates an ability to get code into the mainline tends not to have trouble finding job offers.**”

Who should take this course?

Anyone wants to work on the above problems

Anyone cares about what's going on under the hood

Anyone has to build high-performance systems

Anyone needs to diagnose bugs or security problems

...



About this course

CS 594: Linux Kernel Programming (special topics)

- A little different from a traditional OS course
- Avoid teaching some important but not commonly used OS concepts
 - E.g., boot-loading, enable paging, page table creation, ...
- How to more confidently program inside the Linux kernel

About this course

CS 594: Linux Kernel Programming (special topics)

- A little different from a traditional OS course

Goals

- Understand core subsystems of the Linux kernel in depth
- Design, implement, and modify Linux kernel code and modules for these subsystems
- Test, debug and evaluate the performance of systems software in kernel or user space, using debugging, monitoring and tracing tools


Prerequisite

Required skills:

- C programming and Linux command line (required)
- Basic knowledge of computer architecture, operating systems, algorithms, and data structures (highly recommended)

Textbook

Robert Love, Linux Kernel Development, Addison-Wesley




Other useful resources

- Billimoria, Kaiwan N. Linux Kernel Programming: A comprehensive guide to kernel internals, writing kernel modules, and kernel synchronization. Packt Publishing Ltd, 2021. (highly recommended)
- Billimoria, Kaiwan N. Linux Kernel Debugging: Leverage proven tools and advanced techniques to effectively debug Linux kernels and kernel modules. Packt Publishing Ltd, 2022.
- Bovet, Daniel P., and Marco Cesati. Understanding the Linux Kernel: from I/O ports to process management (3rd Edition). O'Reilly Media, Inc., 2005.
- Corbet, Jonathan, Alessandro Rubini, and Greg Kroah-Hartman. Linux Device Drivers (3rd Edition). O'Reilly Media, Inc., 2005.
- Love, Robert. Linux System Programming: Talking Directly to the Kernel and C Library (2nd Edition). O'Reilly Media, Inc., 2013.

Other useful resources

- Billimoria, Kaiwan N. Linux Kernel Programming: A comprehensive guide to kernel internals, writing kernel modules, and kernel synchronization. Packt Publishing Ltd, 2021. (highly recommended)
- Billimoria, Kaiwan N. Linux Kernel Debugging: Leverage proven tools and advanced techniques to effectively debug Linux kernels and kernel modules. Packt Publishing Ltd, 2021.
- Bovet, Daniel P., and Marco Cesati. Understanding the Linux Kernel: from I/O management to process management (3rd Edition). O'Reilly Media, Inc., 2005.
- Corbet, Jonathan, Alessandro Rubini, and Greg Kroah-Hartman. Linux Device Drivers (3rd Edition). O'Reilly Media, Inc., 2005.
- Love, Robert. Linux System Programming: Talking Directly to the Kernel and C Library (3rd Edition). O'Reilly Media, Inc., 2013.



Communication

Lecture

- Location: TBH 180E
- Time: TR 9:30am – 10:45am
- Recorded video will be uploaded to Blackboard
 - Strongly encourage you to join in person

Office Hours

- Thursday 3:30 PM - 4:30 PM, SEO 1331
- Or by appointment: <https://calendly.com/xgwang9/15-min-office-hour>

Communication

Blackboard

- https://uic.blackboard.com/ultra/courses/_266043_1/cl/outline
- Syllabus, slides, recorded videos, etc.
- Exercise, project submission
- Grades posted

Grading policy

Quizzes and assignments [45%]

- Including paper reading

Midterm exam [15%]

- No final exam

Final research project [30%]

Class participation and discussion [10%]

If you clearly put significant effort into homework, class discussion, the paper(s) to read, and your project, you'll get an A.

Policy for Missed or Late Work

- Late submission (0, 24 hours] will be accepted with a 15% penalty;
- Late submission (24-48 hours] will be accepted with a 30% penalty;
- Late submission beyond 48 hours will not be accepted.

A: ≥ 90 , B: [80,90), C: [70, 80), D: [60, 70), F: < 60

About projects (subject to change)

Tentative topics (maybe only choose 3):

- Adding new system calls
- Kernel module – data structure handling
- Handling page faults from the user-space with `userfaultfd`
- An OS virtualization lab – virtualize a simple random number generator (RNG)

Final research project

- TBD (semester long research project, group of ~2)
- Target high (aim to make a paper submission to SOSP/OSDI/ATC/APSys)

About paper reading

Paper reading with a report

- help you understand a topic better
 - E.g., FlexSC OSDI'10, system call patch processing

Paper reading at the end of the semester (~ 3 weeks)

- everyone read and present 1~2 paper from OSDI / SOSPP / ATC / EuroSys / ASPLOS 2023 / 2024

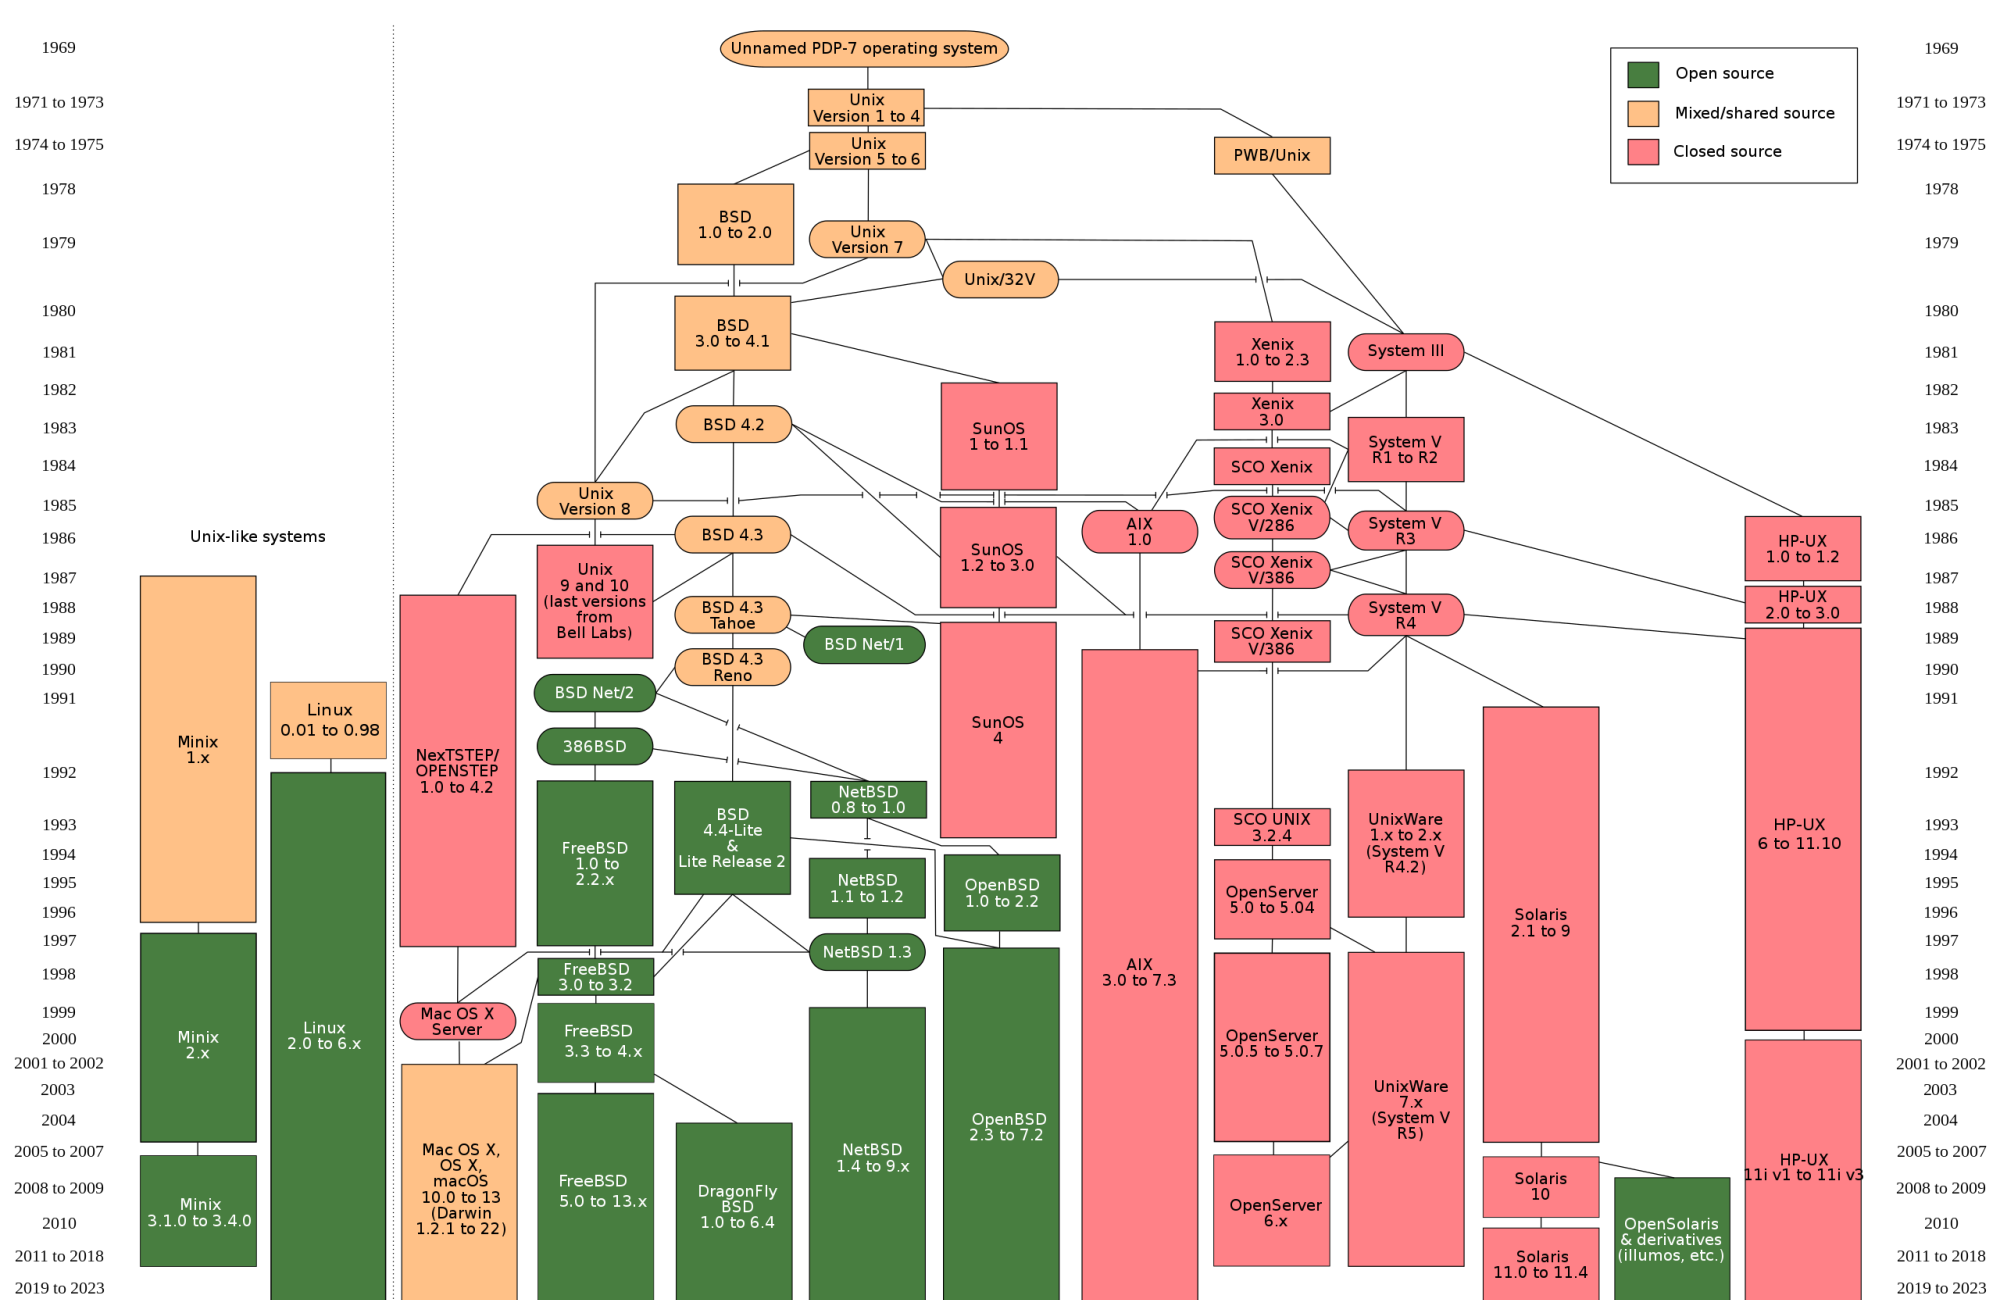
Acknowledgement



Part of the teaching materials are from an ECE-4414/5414G(CS-4224/5264G) course from Virginia Tech provided by [Dr. Changwoo Min](#) and [Dr. Pierre Olivier](#).

Today's agenda

- The history of Linux
- Linux open-source model and community
- High level overview of the Linux kernel



History of UNIX: https://en.wikipedia.org/wiki/History_of_Unix

Beginning of Linux

From: torvalds@klaava.Helsinki.FI (Linus Benedict Torvalds)
Newsgroups: comp.os.minix
Subject: What would you like to see most in minix?
Summary: small poll for my new operating system
Message-ID: <1991Aug25.205708.9541@klaava.Helsinki.FI>
Date: 25 Aug 91 20:57:08 GMT
Organization: University of Helsinki

Hello everybody out there using minix -

I'm doing a (free) operating system (just a hobby, won't be big and professional like gnu) for 386(486) AT clones. This has been brewing since april, and is starting to get ready. I'd like any feedback on things people like/dislike in minix, as my OS resembles it somewhat (same physical layout of the file-system (due to practical reasons) among other things).

I've currently ported bash(1.08) and gcc(1.40), and things seem to work. This implies that I'll get something practical within a few months, and I'd like to know what features most people would want. Any suggestions are welcome, but I won't promise I'll implement them 😊

Linus (torvalds@kruuna.helsinki.fi)

PS. Yes - it's free of any minix code, and it has a multi-threaded fs. It is NOT protable (uses 386 task switching etc), and it probably never will support anything other than AT-harddisks, as that's all I have :-).

Linux history

1991: First apparition, author Linux Torvalds

1992: GPL License, first Linux distribution

1994: v1.0 – Single CPU for i386, then ported to Alpha, MIPS

1996: v2.0 – Symmetric multiprocessing (SMP) support

1999: v2.2 – Big kernel lock removed

2003: v2.6 – Physical address expansion (PAE), new architectures, ...

2011: v3.0 – Incremental release of v2.6

2015: v4.0 – Livepatch

→ today's latest version: <https://kernel.org/>

Linux open-source model



Linux is licensed under GPLv2

“You may copy, distribute and modify the software as long as you track changes/dates in source files. Any modifications to or software including (via compiler) GPL-licensed code must also be made available under the GPL along with build & install instructions.”

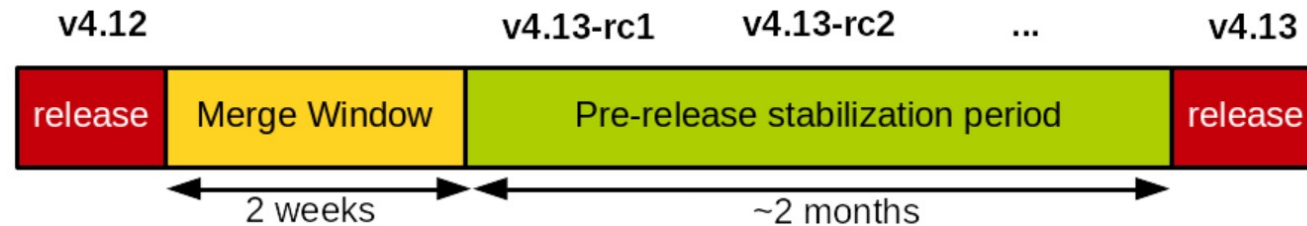
Benefit of open-source model

Given enough eyeballs, all bugs are shallow


Given a large enough tests, almost every problem will be characterized quickly and the fix obvious to someone.



Kernel release cycle

(major).(minor).(stable) → E.g., 6.1.71




- “RC” (Release Candidate) → testing before the mainline release
- Mainline release → maintained by Linus with all new features
- Stable release → additional bug fixes after the mainline kernel release
- Long term support (LTS) for a subset of releases → e.g., 6.1, 5.15, ...

 kernel.org



The Linux Kernel Archives

[About](#)[Contact us](#)[FAQ](#)[Releases](#)[Signatures](#)[Site news](#)



Protocol

Location


[HTTP](#)[https://www.kernel.org/pub/](#)

[GIT](#)[https://git.kernel.org/](#)

[RSYNC](#)[rsync://rsync.kernel.org/pub/](#)

Latest Release

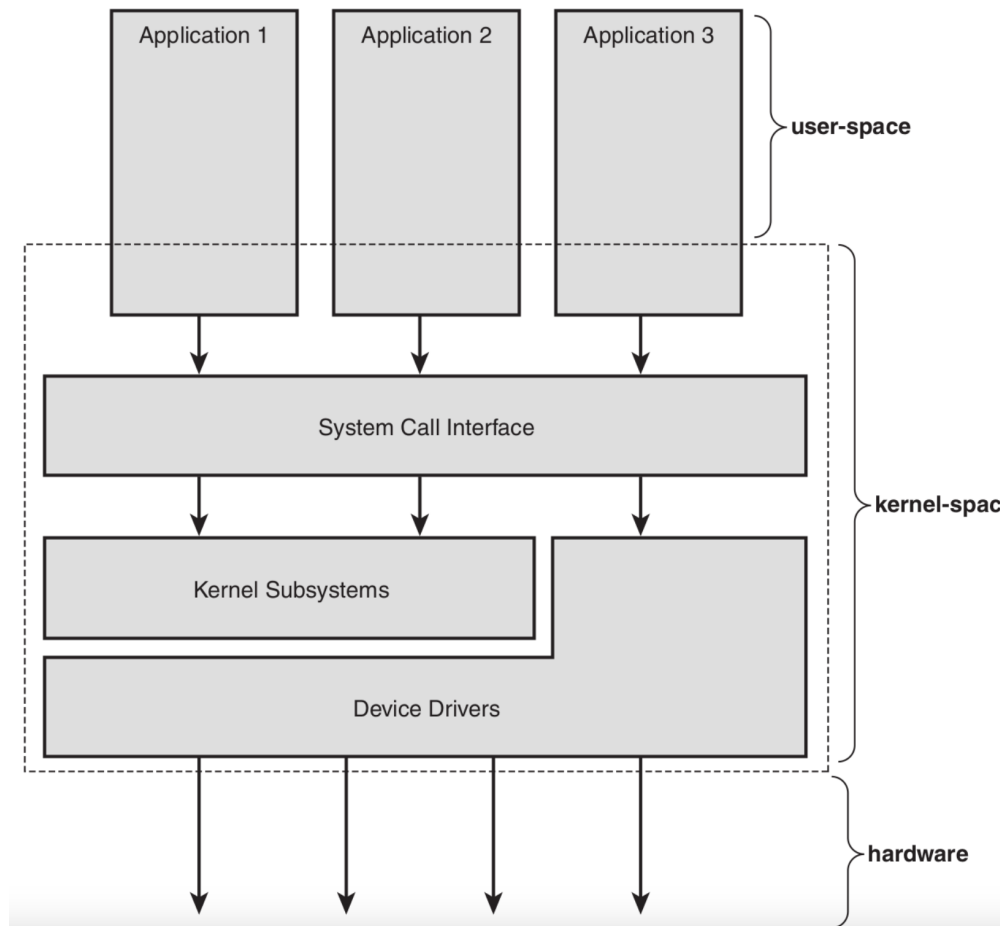
6.6.10



mainline:	6.7-rc8	2023-12-31	[tarball]	[patch]	[inc. patch]	[view diff]	[browse]	
stable:	6.6.10	2024-01-05	[tarball]	[pgp]	[patch]	[inc. patch]	[view diff]	[browse] [changelog]
longterm:	6.1.71	2024-01-05	[tarball]	[pgp]	[patch]	[inc. patch]	[view diff]	[browse] [changelog]
longterm:	5.15.146	2024-01-05	[tarball]	[pgp]	[patch]	[inc. patch]	[view diff]	[browse] [changelog]
longterm:	5.10.206	2024-01-05	[tarball]	[pgp]	[patch]	[inc. patch]	[view diff]	[browse] [changelog]
longterm:	5.4.265	2023-12-20	[tarball]	[pgp]	[patch]	[inc. patch]	[view diff]	[browse] [changelog]
longterm:	4.19.303	2023-12-20	[tarball]	[pgp]	[patch]	[inc. patch]	[view diff]	[browse] [changelog]
longterm:	4.14.334	2023-12-20	[tarball]	[pgp]	[patch]	[inc. patch]	[view diff]	[browse] [changelog]
linux-next:	next-20240105	2024-01-05						[browse]

<https://kernel.org/>

Overview of operating systems



User space v.s. kernel space

A CPU is executing in either of **user space** or **kernel space**

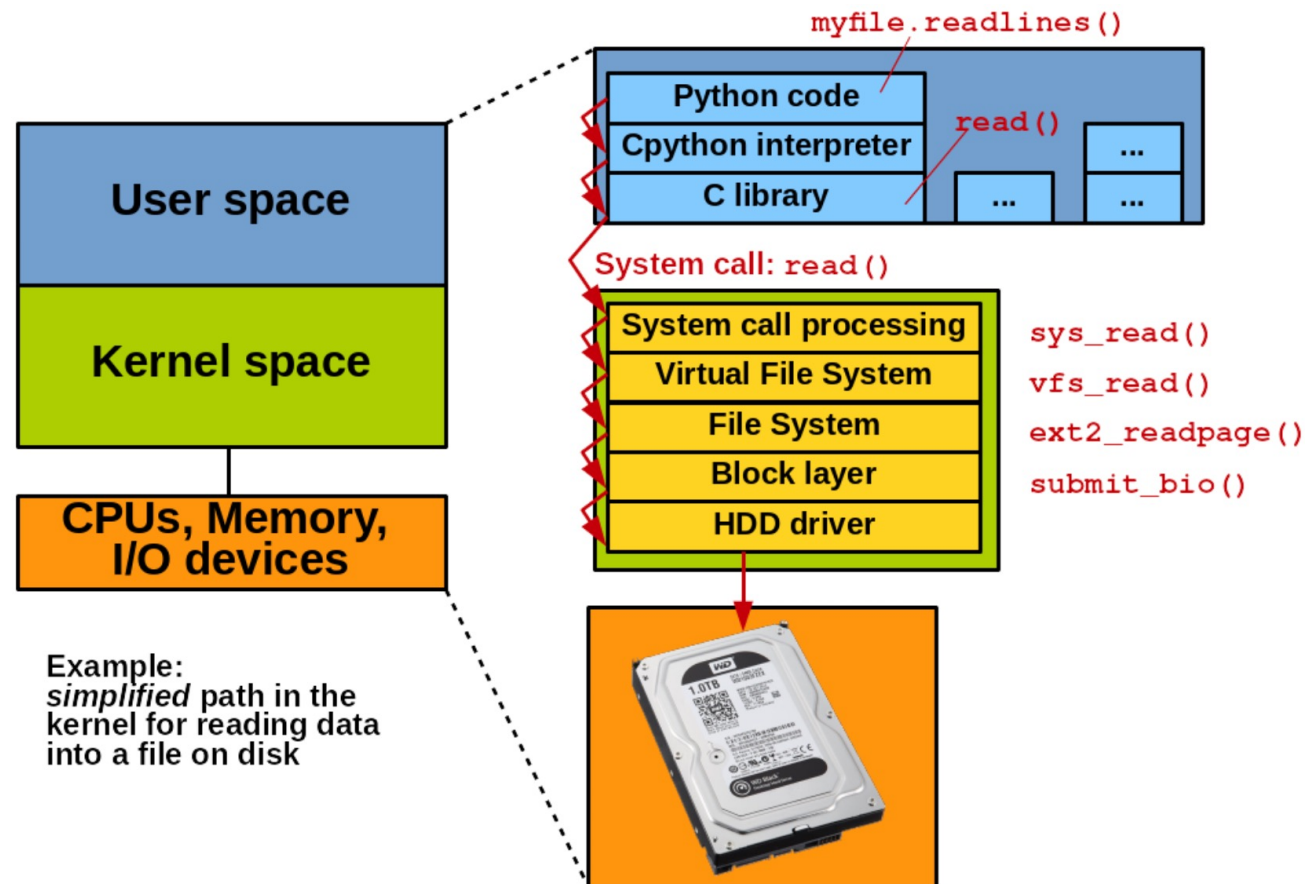
Only the kernel is allowed to perform **privileged operations** such as controlling CPU and IO devices

- E.g., protection rings in x86 architecture
- ring 3: user-space application; ring 0: OS kernel

A user-space application talks to the kernel through the **system call** interface

- E.g., `open()`, `read()`, `write()`, `close()`

User space v.s. kernel space



Linux is a monolithic kernel

A traditional design: all of the OS runs in kernel, privileged mode

- share the same address space

Kernel interface \sim system call interface

Good: easy for subsystems to cooperate

- one cache shared by file system and virtual memory

Bad: leads to bugs, no isolation within kernel

Alternative: micro-kernel design

Many OS services run as ordinary user programs

- e.g., file system in a file server

Kernel implements minimal mechanism to run services in user space

- IPC, virtual memory, threads

Kernel interface != system call interface

- applications talk to servers via IPCs

Good: more isolation

Bad: IPCs may be slow

Debate

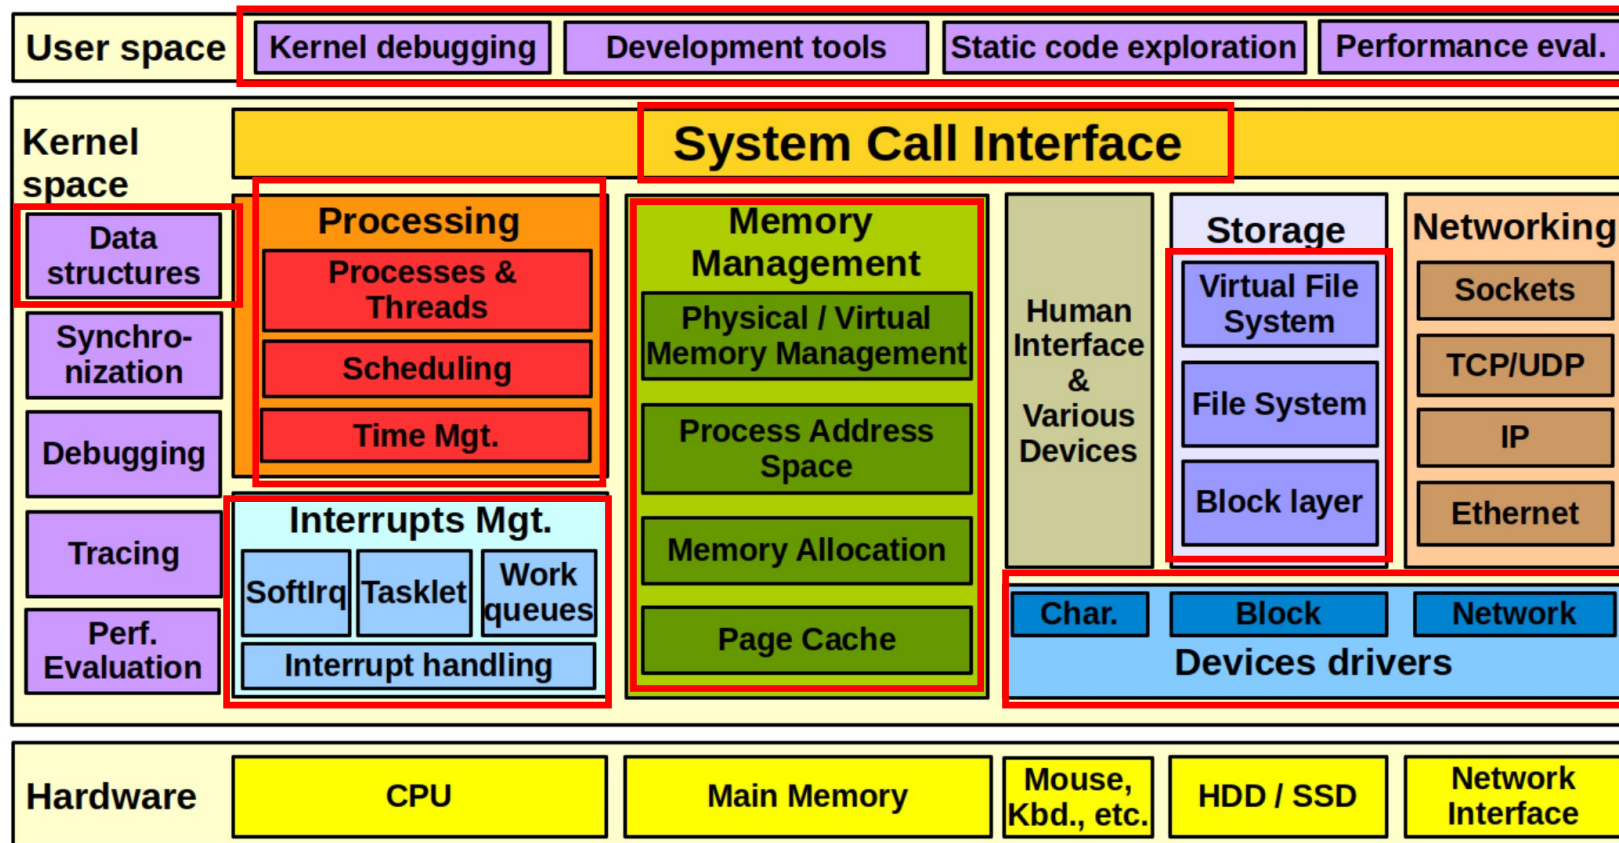


Tanenbaum-Torvalds debate

Most real-world kernels are mixed: Linux, OS X, Windows

- e.g., X Window system

Kernel & course map



Set up course environment

VirtualBox to run Linux VM

- Recommended setting
 - disk \geq 64GB, RAM \geq 4GB, # CPU \geq 2
- Add port forwarding rule
 - protocol: TCP, host IP: 127.0.0.1
 - host port: 2222, guest port: 22

Cloudlab

- <https://cloudlab.us/>

Set up course environment

Ubuntu 22.04 server for Linux distribution

- Add your account as a sudo user

```
sudo usermod -aG sudo xiaoguang
```

- [Tricks to avoid typing password every time](#)

```
sudo visudo
```

```
xiaoguang ALL=(ALL) NOPASSWD:ALL
```

Linux kernel: v6.1 released on Dec 12th, 2022

Next steps

Finish to set up course environment

- Bring your laptop

Take the Readiness Exercise (hw1)

- Due Friday!

If you are not familiar with Linux commands, learn followings:

- vim, ssh, scp, tmux, and [more](#)

Download the latest Linux kernel source inside your Linux VM

- `$ git clone https://github.com/torvalds/linux.git`

Next lecture



Build and explore Linux kernel

Survey

