

Interrupt Handler: Bottom Half

Xiaoguang Wang

Homework

hw3 grading is out

paper reading due yesterday (Feb/12th)

hw4 & hw5: kernel module

- use the **ubuntu cloud image** if you have difficulty to access a Linux machine
- Instructions on blackboard discussion (Environment Settings)

Recap: Interrupt controller



Interrupts are electrical signals multiplexed by the **interrupt controller**

- Sent on a specific pin of the CPU

Once an interrupt is received, a dedicated function is executed

- **Interrupt handler**

The kernel/user space can be interrupted at (nearly) any time to process an interrupt

Recap: Advanced PIC (APIC, I/O APIC)

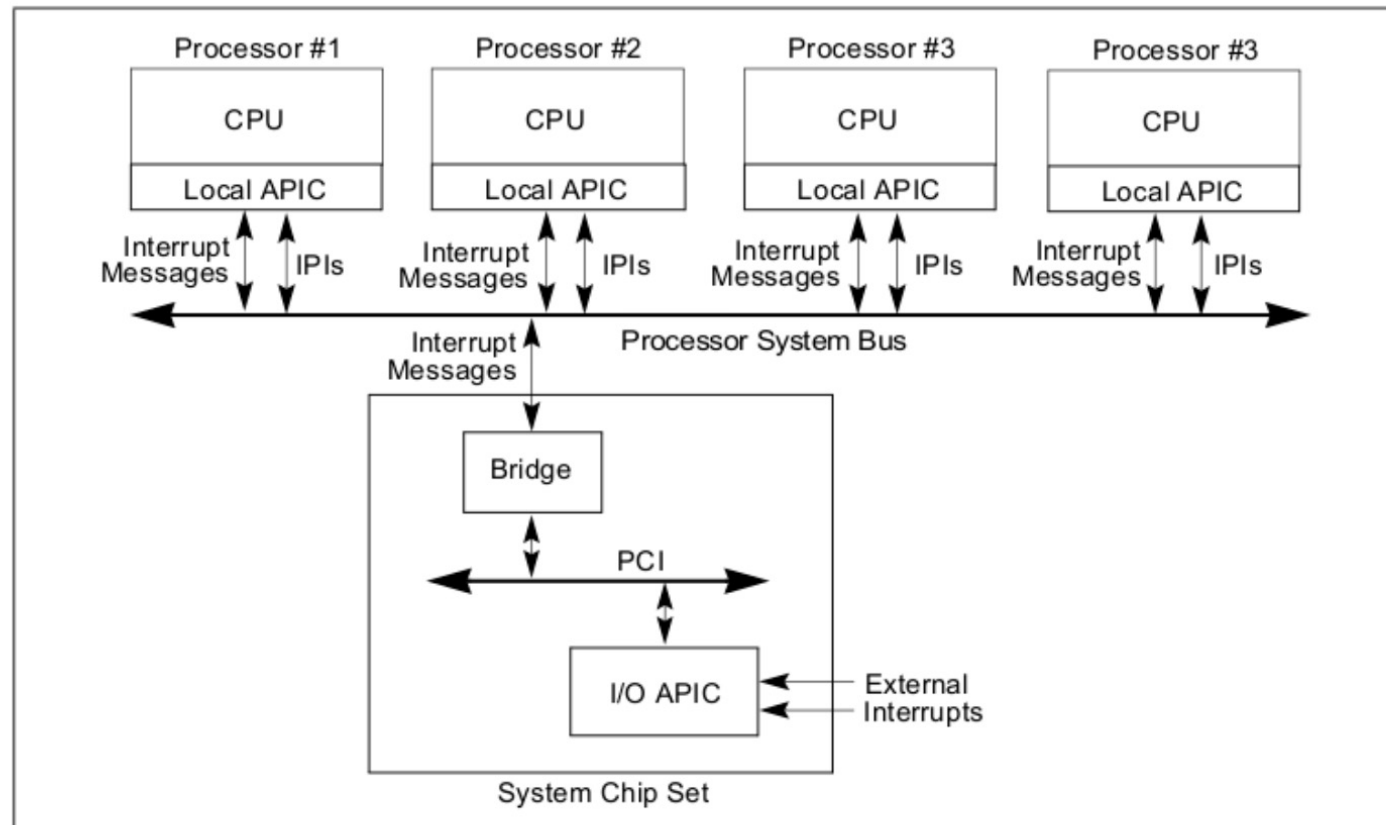
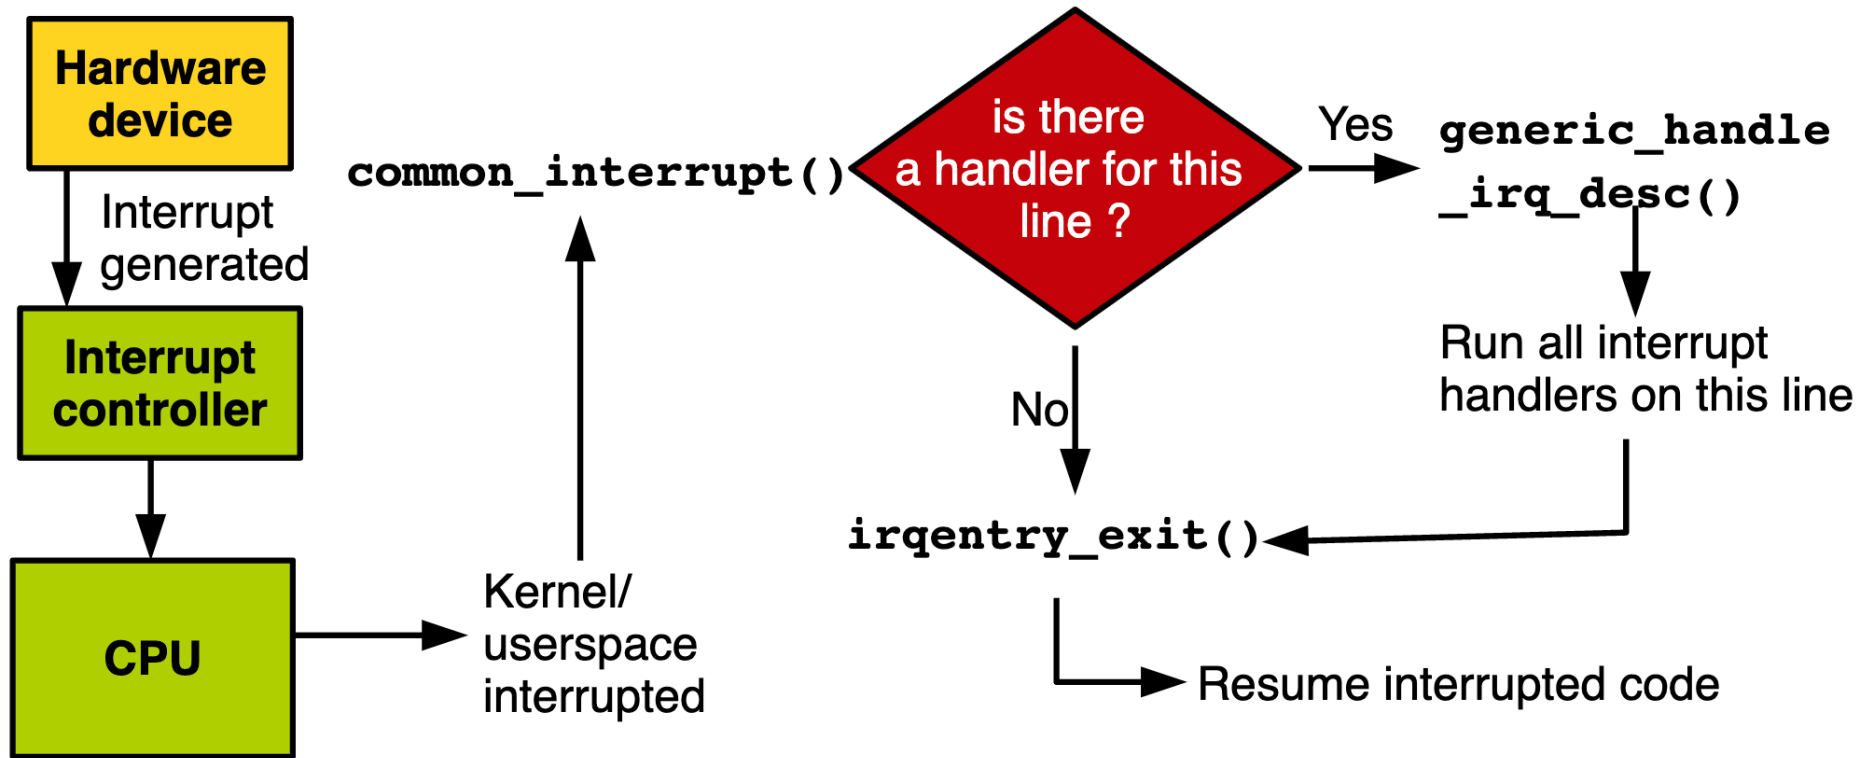


Figure 10-2. Local APICs and I/O APIC When Intel Xeon Processors Are Used in Multiple-Processor Systems

Recap: Interrupt handling in Linux



Today: interrupt handler

Top-halves (interrupt handlers) must run as quickly as possible

- They are interrupting other kernel/user code
- They are often timing-critical because they deal with hardware.
- They run in interrupt context: **they cannot block**
- **One or all interrupt lines are disabled**

Defer the less critical part of interrupt processing to a **bottom-half**

The history of bottom halves

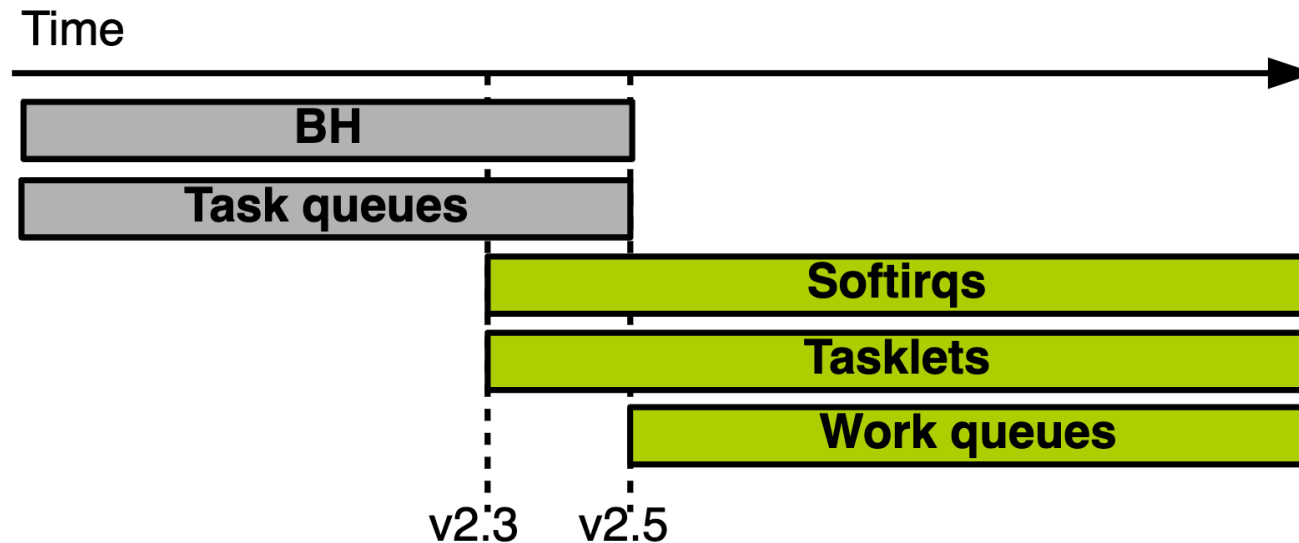
“Top-half” and “bottom-half” are generic terms not specific to Linux
old “Bottom-Half” (BH) mechanism

- a statistically created list of 32 bottom halves

Task queues: queues of function pointers

- still too inflexible
- not lightweight enough for performance-critical subsystems (e.g., networking)


The history of bottom halves



BH → Softirq, tasklet

Task queue → work queue

Today's bottom halves in Linux



All bottom-half mechanisms **run with all interrupts enabled**

Softirqs and **tasklets** run in an interrupt context

- Softirq is rarely used directly
- Tasklet is a simple and easy-to-use softirq (built on softirq)

Work queues run in a process context

- They can block and go to sleep

Softirq

```
/* kernel/softirq.c */
/* Softirq is statically allocated at compile time */
static struct softirq_action softirq_vec[NR_SOFTIRQS]; /* softirq vector */
/* include/linux/interrupt.h */
enum {
    HI_SOFTIRQ=0,      /* [highest priority] high-priority tasklet */
    TIMER_SOFTIRQ,     /* timer */
    NET_TX_SOFTIRQ,     /* send network packets */
    NET_RX_SOFTIRQ,     /* receive network packets */
    BLOCK_SOFTIRQ,      /* block devices */
    IRQ_POLL_SOFTIRQ,   /* interrupt-poll handling for block device */
    TASKLET_SOFTIRQ,    /* normal priority tasklet */
    SCHED_SOFTIRQ,      /* scheduler */
    HRTIMER_SOFTIRQ,    /* unused */
    RCU_SOFTIRQ,        /* [lowest priority] RCU locking */

    NR_SOFTIRQS         /* the number of defined softirq (< 32) */
};
struct softirq_action {
    void (*action)(struct softirq_action *); /* softirq handler */
};
```

Execute softirq

Raise a softirq (by the h/w interrupt handler)

- Mark the execution of a particular softirq is needed
- Usually, a top-half marks its softirq for execution before returning

Pending softirqs are checked and executed in following places:

- In the return from hardware interrupt code path
- In the ksoftirqd kernel thread
- In any code that explicitly checks for and executes pending softirqs

Execute softirq

Go over the softirq vector and executes the pending softirq handler

```
/* kernel/softirq.c */
/* do_softirq() calls __do_softirq() */
void __do_softirq(void) /* much simplified version for explanation */
{
    u32 pending;
    pending = local_softirq_pending(); /* 32-bit flags for pending softirq */
    if (pending) {
        struct softirq_action *h;

        set_softirq_pending(0); /* reset the pending bitmask */
        h = softirq_vec;
        do {
            if (pending & 1)
                h->action(h); /* execute the handler of the pending softirq */
            h++;
            pending >>= 1;
        } while (pending);
    }
}
```

Use softirq: assign an index

```
enum {  
    HI_SOFTIRQ=0,           /* include/linux/interrupt.h */  
    TIMER_SOFTIRQ,         /* [highest priority] high-priority tasklet */  
    NET_TX_SOFTIRQ,        /* timer */  
    NET_RX_SOFTIRQ,        /* send network packets */  
    BLOCK_SOFTIRQ,         /* receive network packets */  
    IRQ_POLL_SOFTIRQ,      /* block devices */  
    TASKLET_SOFTIRQ,       /* interrupt-poll handling for block device */  
    SCHED_SOFTIRQ,         /* normal priority tasklet */  
    HRTIMER_SOFTIRQ,       /* scheduler */  
    RCU_SOFTIRQ,           /* unused */  
    YOUR_NEW_SOFTIRQ,      /* [lowest priority] RCU locking */  
    NR_SOFTIRQS            /* TODO: add your new softirq index, but not recommended */  
}
```

Use softirq: register a handler

```
/* kernel/softirq.c */  
/* register a softirq handler for nr */  
void open_softirq(int nr, void (*action)(struct softirq_action *))  
{  
    softirq_vec[nr].action = action;  
}
```

```
/* net/core/dev.c */  
static int __init net_dev_init(void)  
{  
    /* ... */  
    /* register softirq handler to send messages */  
    open_softirq(NET_TX_SOFTIRQ, net_tx_action);  
  
    /* register softirq handler to receive messages */  
    open_softirq(NET_RX_SOFTIRQ, net_rx_action);  
    /* ... */  
}
```

handler for network message send

Use softirq: register a handler

Run with interrupts enabled and cannot sleep

The key advantage of softirq over tasklet is scalability

- If the same softirq is raised again while it is executing, another processor can run it simultaneously

This means that any shared data needs proper locking

- To avoid locking, most softirq handlers resort to per-processor data (data unique to each processor and thus not requiring locking)

Use softirq: raise a softirq

Softirqs are most often raised from within interrupt handlers (i.e., top halves)

- The interrupt handler performs the basic hardware-related work, raises the softirq, and then exits

```
/* include/linux/interrupt.h */  
/* Disable interrupt and raise a softirq */  
extern void raise_softirq(unsigned int nr);  
  
/* Raise a softirq. Interrupt must already be off. */  
extern void raise_softirq_irqoff(unsigned int nr);  
  
/* kernel/time/timer.c */  
run_local_timers() --> raise_softirq(TIMER_SOFTIRQ);
```




Tasklet

Built on top of softirqs

- `HI_SOFTIRQ`: high priority tasklet
- `TASKLET_SOFTIRQ`: normal priority tasklet

Running in an interrupt context (i.e., cannot sleep)

- Like softirq, **all interrupts** are **enabled**

Restricted concurrency than softirq

- The same tasklet cannot run concurrently

tasklet_struct

```
/* include/linux/interrupt.h */
struct tasklet_struct {
    struct tasklet_struct *next;    /* next tasklet in the list */
    unsigned long state;           /* state of a tasklet: scheduled, running */
    atomic_t count;                /* disable counter: != 0 cannot run */
    void (*func)(unsigned long); /* tasklet handler function */
    unsigned long data;            /* argument of the tasklet function */
};
```

Schedule a tasklet

Scheduled tasklets are stored in two **per-processor linked list**:

- tasklet_vec, tasklet_hi_vec

```
/* kernel/softirq.c */
struct tasklet_head {
    struct tasklet_struct *head;
    struct tasklet_struct **tail;
};

/* regular tasklet */
static DEFINE_PER_CPU(struct tasklet_head, tasklet_vec);
/* high-priority tasklet */
static DEFINE_PER_CPU(struct tasklet_head, tasklet_hi_vec);
```

Schedule a tasklet

```
/* include/linux/interrupt.h, kernel/softirq.c */  
  
/* Schedule a regular tasklet; For high-priority tasklet, use tasklet_hi_schedule() */  
static inline void tasklet_schedule(struct tasklet_struct *t)  
{  
    if (!test_and_set_bit(TASKLET_STATE_SCHED, &t->state))  
        __tasklet_schedule(t);  
}  
  
void __tasklet_schedule(struct tasklet_struct *t)  
{  
    __tasklet_schedule_common(t, &tasklet_vec,  
                             TASKLET_SOFTIRQ);  
}  
EXPORT_SYMBOL(__tasklet_schedule);
```

Schedule a tasklet

```
static void __tasklet_schedule_common(struct tasklet_struct *t,
                                     struct tasklet_head __percpu *headp,
                                     unsigned int softirq_nr)
{
    struct tasklet_head *head;
    unsigned long flags;

    local_irq_save(flags);           /* disable interrupt */
    head = this_cpu_ptr(headp);      /* Append this tasklet at the end of list */
    t->next = NULL;
    *head->tail = t;
    head->tail = &(t->next);
    raise_softirq_irqoff(softirq_nr); /* tasklet is a softirq */
    local_irq_restore(flags);        /* enable interrupt */
}
```

Tasklet softirq handlers

```
/* kernel/softirq.c */
void __init softirq_init(void)
{
    /* ... */
    /* Tasklet softirq handlers are registered at initializing softirq */
    open_softirq(TASKLET_SOFTIRQ, tasklet_action);
    open_softirq(HI_SOFTIRQ, tasklet_hi_action);
}

static __latent_entropy void tasklet_action_comment(struct softirq_action *a,
                                                    struct tasklet_head *tl_head,
                                                    unsigned int softirq_nr)
{
    struct tasklet_struct *list;

    /* Clear the list for this processor by setting it equal to NULL */
    local_irq_disable();
    list = tl_head->head;
    tl_head->head = NULL;
    tl_head->tail = &tl_head->head;
    local_irq_enable();
}
```

Tasklet softirq handlers (cont'd)

```
/* For all tasklets in the list */
while (list) {
    struct tasklet_struct *t = list;
    list = list->next;
    /* If a tasklet is not processing and it is enabled */
    if (tasklet_trylock(t) && !atomic_read(&t->count)) {
        /* and it is not running */
        if (!test_and_clear_bit(TASKLET_STATE_SCHED, &t->state))
            BUG();
        /* then execute the associate tasklet handler */
        t->func(t->data);
        tasklet_unlock(t);
        continue;
    }
    tasklet_unlock(t);
}
local_irq_disable();
t->next = NULL;
*tl_head->tail = t;
tl_head->tail = &t->next;
raise_softirq_irqoff(softirq_nr);
local_irq_enable();
}
```


Use tasklet: declaring a tasklet

```
/* include/linux/interrupt.h */

/* Static declaration of a tasklet with initially enabled */
#define DECLARE_TASKLET(name, _callback) \
struct tasklet_struct name = { .count = ATOMIC_INIT(0), /* disable counter */ \
                               .callback = _callback, \
                               .use_callback = true, }

/* Static declaration of a tasklet with initially disabled */
#define DECLARE_TASKLET_DISABLED(name, _callback) \
struct tasklet_struct name = { .count = ATOMIC_INIT(1), /* disable counter */ \
                               .callback = _callback, \
                               .use_callback = true, }

/* Dynamic initialization of a tasklet */
extern void tasklet_init(struct tasklet_struct *t,
                        void (*func)(unsigned long), unsigned long data);
```

Use tasklet: tasklet handler

Run with interrupts enabled and cannot sleep

- If your tasklet shared data with an interrupt handler, you must task precautions (e.g., disable interrupt or obtain a lock)

Two of the same tasklets never run concurrently

- Because `tasklet_action()` checks `TASKLET_STATE_RUN`

But two different tasklets can run at the same time on two different processors

Schedule a tasklet

```
/* include/linux/interrupt.h */  
  
void tasklet_schedule(struct tasklet_struct *t);  
void tasklet_hi_schedule(struct tasklet_struct *t);  
  
/* Disable a tasklet by increasing the disable counter */  
void tasklet_disable(struct tasklet_struct *t)  
{  
    tasklet_disable_nosync(t);  
    tasklet_unlock_wait(t);           /* and wait until the tasklet finishes */  
    smp_mb();  
}  
  
void tasklet_enable(struct tasklet_struct *t)  
{  
    smp_mb__before_atomic();  
    atomic_dec(&t->count);  
}
```

Tasklet example

Blackboard [lec11-bottom-half.tar.gz](#)

ksoftirqd

Per-processor kernel thread that aids processing softirqs (kernel softirq daemon)

If the number of softirqs grows excessive, the kernel wakes up ksoftirqd with normal priority (nice 0)

- No starvation of user-space application
- Running a softirq has the normal priority (nice 0)

```
► ps ax -eo pid,nice,stat,cmd | grep ksoftirqd
    13      0 S      [ksoftirqd/0]
    22      0 S      [ksoftirqd/1]
    28      0 S      [ksoftirqd/2]
```



Work queues

Work queues defer work into a kernel thread

- Always runs in process context
- Thus, work queues are schedulable and can therefore sleep

By default, per-cpu kernel thread is created, kworker/n

- You can create additional per-CPU worker thread, if needed

Workqueue users can also create their own threads for better performance and lighten the load on default threads

```
▶ ps ax -eo pid,nice,stat,cmd | grep kworker
    8 -20 I<    [kworker/0:0H-events_highpri]
   17   0 I     [kworker/0:1-events]
   24 -20 I<    [kworker/1:0H-events_highpri]
```

Work queue: data structure

```
/* kernel/workqueue.c */
struct worker_pool {
    spinlock_t lock;           /* the pool lock */
    int cpu;                   /* l: the associated cpu */
    int node;                  /* l: the associated node ID */
    int id;                    /* l: pool ID */
    unsigned int flags;        /* X: flags */
    struct list_head worklist; /* L: list of pending works */
    int nr_workers;            /* L: total number of workers */
    /* ... */
};
```

```
/* include/workqueue.h */
struct work_struct {
    atomic_long_t data;
    struct list_head entry;
    work_func_t func;
};
```


Work queue: work thread

Worker threads execute the `worker_thread()` function

Infinite loop doing the following:

1. Check if there is some work to do in the current pool
2. If so, execute all the `work_struct` objects pending in the pool `worklist` by calling `process_scheduled_works()`
 - Call the `work_struct` function pointer `func`
 - `work_struct` objects removed
3. Go to sleep until a new work is inserted in the work queue

Work queue: create work

```
/* include/linux/workqueue.h */
```

```
/* Statically creating a work */
```

```
DECLARE_WORK(work, handler_func);
```

```
/* Dynamically creating a work at runtime */
```

```
INIT_WORK(work_ptr, handler_func);
```

```
/* Work handler prototype */
```

```
typedef void (*work_func_t)(struct work_struct *work);
```

```
/* Create/destroy a new work queue in addition to the default queue */
```

```
struct workqueue_struct *create_workqueue(char *name);
```

```
void destroy_workqueue(struct workqueue_struct *wq);
```



Work queue: schedule work

/ Put work task in global workqueue (kworker/n) */*

```
bool schedule_work(struct work_struct *work);
```

```
bool schedule_work_on(int cpu,  
                      struct work_struct *work); /* on the specified CPU */
```

/ Queue work on a specified workqueue */*

```
bool queue_work(struct workqueue_struct *wq, struct work_struct *work);
```

```
bool queue_work_on(int cpu, struct workqueue_struct *wq,  
                  struct work_struct *work); /* on the specified CPU */
```

Work queue: finish work

/ Flush a specific work_struct */*

int flush_work(**struct** work_struct *work);

/ Flush a specific workqueue: */*

void flush_workqueue(**struct** workqueue_struct *);

/ Flush the default workqueue (kworkers): */*

void flush_scheduled_work(**void**);

/ Cancel the work */*

void flush_workqueue(**struct** workqueue_struct *wq);

/ Check if a work is pending */*

work_pending(**struct** work_struct *work);

Work queue example

Blackboard [lec11-bottom-half.tar.gz](#)

Choose the right bottom-half

Bottom half	Context	Inherent serialization
Softirq	Interrupt	None
Tasklet	Interrupt	Against the same tasklet
Work queue	Process	None

All these generally run with interrupts enabled

If there is a **shared data** with **an interrupt handler (top-half)**, need to disable interrupts or use locks

Disable softirq and tasklet

/ Disable softirq and tasklet processing on the local processor */*

```
void local_bh_disable();
```

/ Enable softirq and tasklet processing on the local processor */*

```
void local_bh_enable();
```

These calls can be nested

- Only the final call to `local_bh_enable()` enables bottom halves

These calls do not disable workqueues processing

Next steps

hw6: CPU profiler (part 1)

- due Feb 16th

Reading assignment:

- [Optimizing Storage Performance with Calibrated Interrupts](#), OSDI'21
- due Feb 19th

Further readings

LKD3: Chapter 8: Bottom Halves and Deferring Work