

# Isolation and System Calls

Xiaoguang Wang

# Summary of the last lecture

---

Get, build, and explore the Linux kernel

- `git`, `tig`, `make`, `make modules`, `make modules_install`, `make install`, `vim`, `emacs`, `LXR`, `cscope`, `ctags`, `tmux`

Don't try to master them at once. Instead, gradually get used to them.

# Questions

---

Other vim tools?

- vim-plug

```
git clone --depth=1 https://github.com/xgwan9/.vim.git ~/.vim
```

```
cd ~/.vim
```

```
sh install.sh
```

How do I read Linux kernel code?

# How to read kernel code

---

## E.g., ext4 file system

1. General understanding of OS file systems ← any OS textbook
2. File system in Linux kernel
3. Check kernel Documentation and Ext4 on-disk layout
4. Read the ext4 kernel code
  - module by module; start from a system call (e.g., how `sys_write()` is implemented?)
5. Search LWN to check the latest changes → E.g., ext4 encryption support

# How to read kernel code

Use function tracer

- ftrace: function tracer framework
- perf tools: ftrace front end
  - Try kernel/funcgraph
- Try bpftrace!

```
cs594-s24/kernel/perf-tools master ✓
▶ sudo ./kernel/funcgraph -HtPp 22842 vfs_read
Tracing "vfs_read" for PID 22842... Ctrl-C to end.
# tracer: function_graph
#
#      TIME      CPU  TASK/PID      DURATION
#      |         |   |         |         |
2664.970859 | 2)  zsh-22842 |         |
2664.970862 | 2)  zsh-22842 |         |
2664.970863 | 2)  zsh-22842 | 0.431 us |
2664.970863 | 2)  zsh-22842 | 1.450 us |
2664.970865 | 2)  zsh-22842 | 0.525 us |
2664.970866 | 2)  zsh-22842 |         |
2664.970866 | 2)  zsh-22842 |         |
2664.970866 | 2)  zsh-22842 |         |
2664.970867 | 2)  zsh-22842 | 0.288 us |
2664.970867 | 2)  zsh-22842 |         |
2664.970867 | 2)  zsh-22842 |         |
2664.970868 | 2)  zsh-22842 |         |
2664.970868 | 2)  zsh-22842 |         |
2664.970869 | 2)  zsh-22842 |         |
2664.970869 | 2)  zsh-22842 | 0.484 us |
2664.970870 | 2)  zsh-22842 |         |
2664.970870 | 2)  zsh-22842 | 0.257 us |
```

```
FUNCTION CALLS
| | | |
finish_task_switch.isra.0() {
  raw_spin_rq_unlock() {
    _raw_spin_unlock();
  }
  irq_enter_rcu();
  __sysvec_irq_work() {
    __wake_up() {
      __wake_up_common_lock() {
        _raw_spin_lock_irqsave();
        __wake_up_common() {
          autoremove_wake_function() {
            default_wake_function() {
              try_to_wake_up() {
                _raw_spin_lock_irqsave();
                select_task_rq_fair() {
                  __rcu_read_lock();
```

# Browse/navigate kernel code

---

```
$ make cscope tags -j2
# cscope - build cscope database; tags - build ctag database
$ vim
# :tag <symbol>          # search symbol definition
# :cs find s <symbol>    # find uses of symbol
# Ctrl - ]              # search symbol definition at the cursor
# Ctrl - t              # return to the previous cursor point
```

# Today: isolation and system calls



How to isolate user applications from the kernel?

How to safely access the kernel from user application?

How does the Linux system call work?

# The unit of isolation: "process"

---

Prevent process X from accessing or spying on process Y

- e.g., memory, address space, FDs, cpu, etc.

Prevent a process from maliciously accessing the operating system itself

- e.g, a buggy or malicious program

How to isolate a process from kernel?



# Isolation mechanisms in OS

---

User/kernel mode flag (a.k.a., rings)

Address spaces (later)

Time-slicing (later)

System call interface

# Hardware isolation in x86

Ring 0: most privileged CPU mode – kernel

Ring 3: most unprivileged CPU mode -- user

**Q: What's the meaning of "rings" here?**

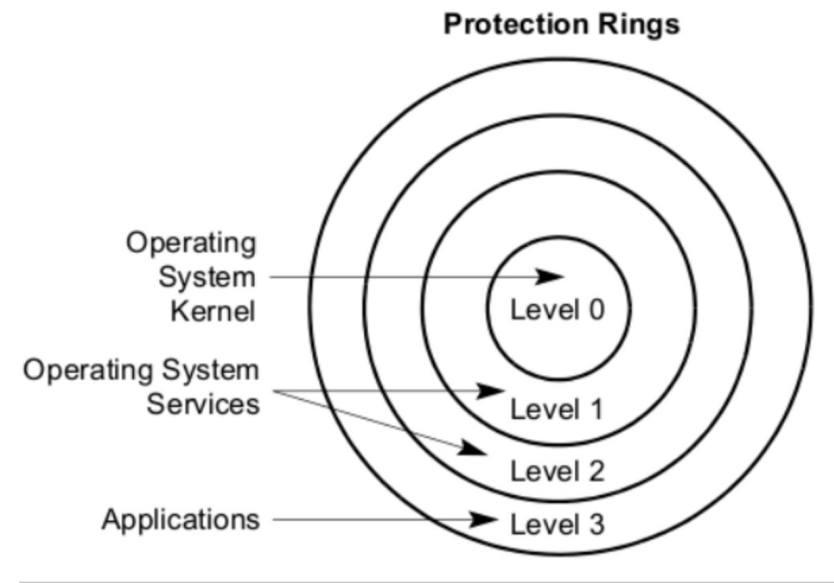


Figure 5-3. Protection Rings

# Segmentation in x86\_64

Logical address:

- segment base + offset

Linear address:

- via page tables → physical address

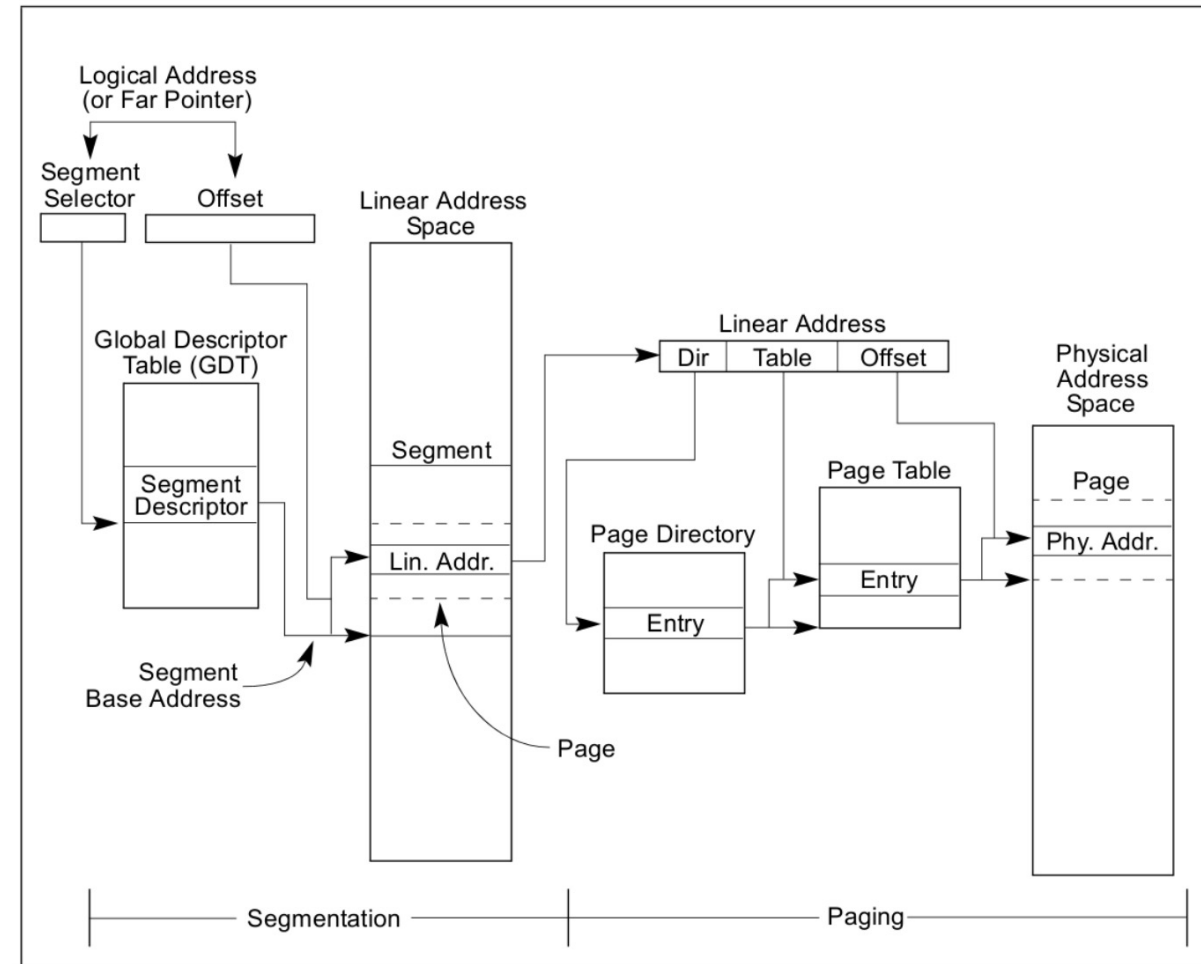



Figure 3-1. Segmentation and Paging

# Segmentation in x86\_64

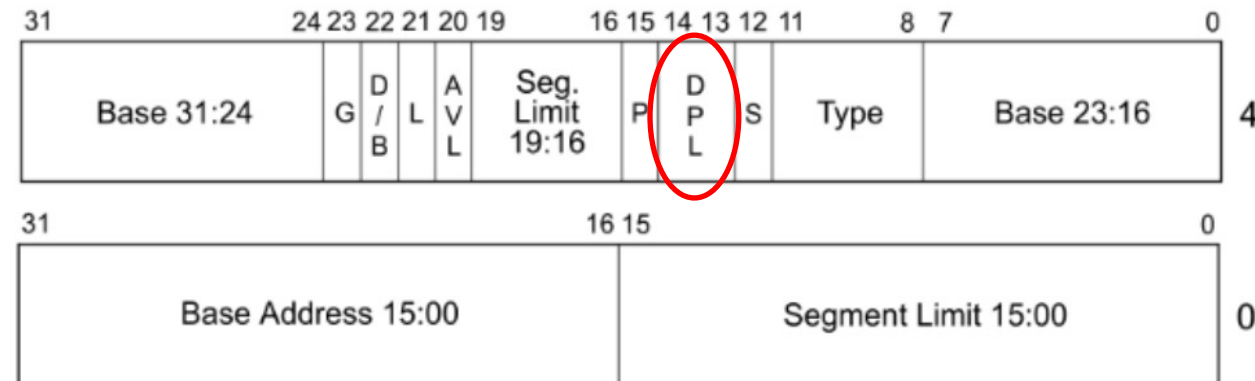
Protected/long mode:

- 16-bit segment registers store a selector
  - %cs, %ds, %ss, %es, %fs, %gs
- a selector contains an index to a segment descriptor table

Segment selector



Segment descriptor (in GDT or LDT)



- L — 64-bit code segment (IA-32e mode only)
- AVL — Available for use by system software
- BASE — Segment base address
- D/B — Default operation size (0 = 16-bit segment; 1 = 32-bit segment)
- DPL — Descriptor privilege level
- G — Granularity
- LIMIT — Segment Limit
- P — Segment present
- S — Descriptor type (0 = system; 1 = code or data)
- TYPE — Segment type

# Segmentation in x86\_64

Protected/long mode:

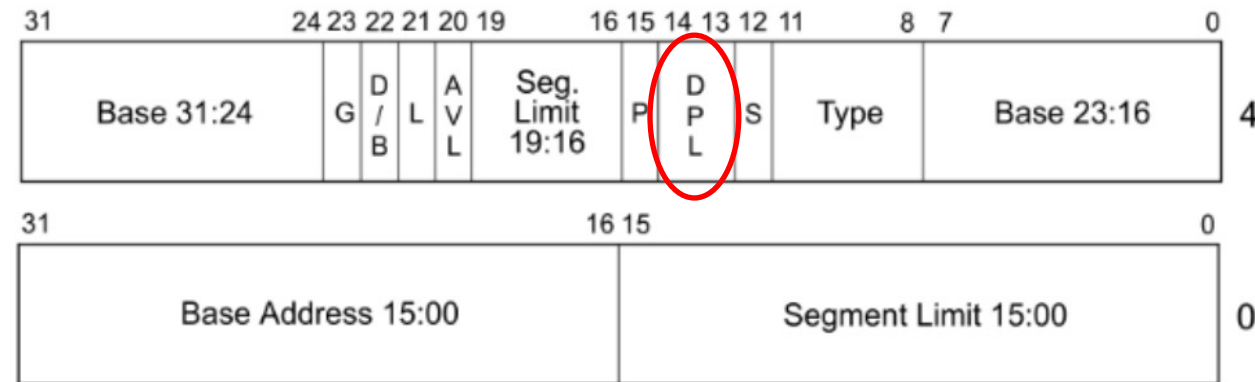
- 16-bit segment registers store a selector
  - %cs, %ds, %ss, %es, %fs, %gs
- a selector contains an index to a segment descriptor table

Segment selector



CS	0x33
SS	0x2b
ds	0x0
es	0x0

Segment descriptor (in GDT or LDT)



L — 64-bit code segment (IA-32e mode only)

AVL — Available for use by system software

BASE — Segment base address

operation size (0 = 16-bit segment; 1 = 32-bit segment)

for privilege level

ity

Limit

nt present

Descriptor type (0 = system; 1 = code or data)

TYPE — Segment type

# Privilege levels of a segment

DPL (descriptor privilege level)

- the privilege level of a segment

CPL (current privilege level)

- the privilege level of currently executing program
- bits 0:1 in the `%cs` register

RPL (requested privilege level)

- an override privilege level that is assigned to a segment selector
- a segment selector is a part (16-bit) of segment registers (e.g., `%ds`, `%fs`), which is an index of a segment descriptor and RPL

# How is isolation enforced in x86?

Access is granted if  $\max(\text{CPL}, \text{RPL}) \leq \text{DPL}$  (x86 segment)

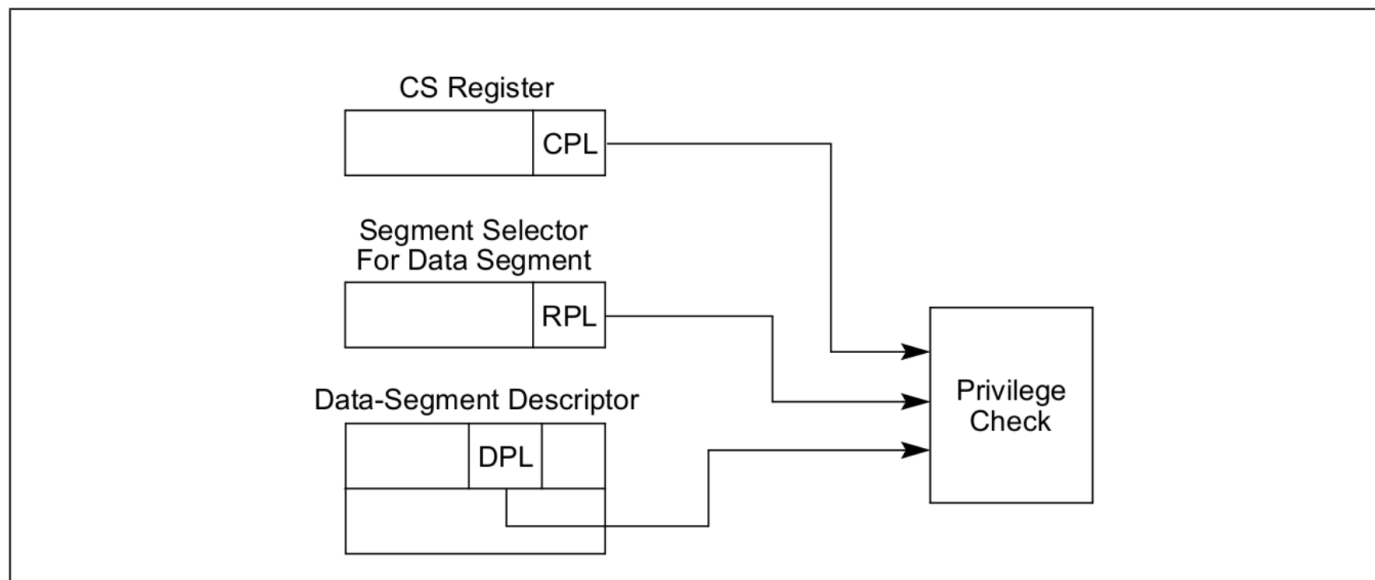


Figure 5-4. Privilege Check for Data Access

# What does "ring 0" protect?

---

Protects everything relevant to isolation

- Writes to %cs (to defend CPL)
- I/O port access
- Control register accesses (eflags, %cr3, ...)



# How to switch between rings?

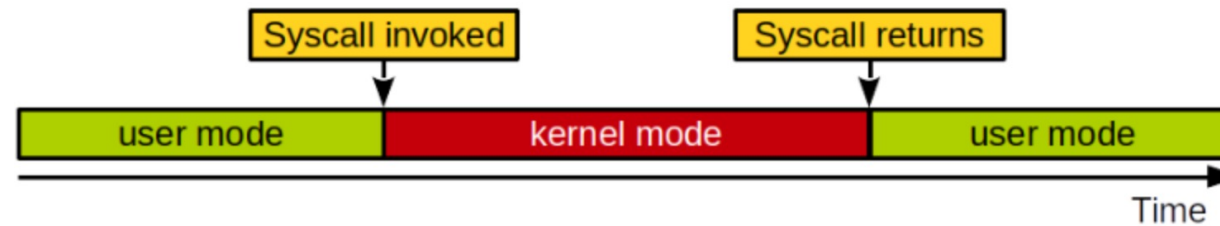
---

Controlled transfer: system calls

- `int 0x80`, `sysenter` or `syscall` instructions set CPL to 0; change to `KERNEL_CS` and `KERNEL_DS` segments
- set CPL to 3 before going back to user space; change to `USER_CS` and `USER_DS` segments

**Q:** How to systematically manage this interface?

# System calls



One way (the only way) for user-space application to enter the kernel to request OS services

- A layer between the hardware and the processes
- An abstract hardware interface for user-space
- Ensure system security and stability

# Examples of system calls

---

- Process management/scheduling: `fork`, `exit`, `execve`, `nice`, `{get|set}priority`, `{get|set}pid`
- Memory management: `brk`, `mmap`
- File system: `open`, `read`, `write`, `lseek`, `stat`
- Inter-Process Communication: `pipe`, `shmget`
- Time management: `{get|set}timeofday`
- Others: `{get|set}uid`, `connect`

**Q: Where are system call implementations in Linux kernel?**

# Syscall table and identifier

The syscall table for the x86\_64 architecture

- `arch/x86/entry/syscalls/syscall_64.tbl`

Syscall ID: unique integer (sequentially assigned)

```
vi arch/x86/entry/syscalls/syscall_64.tbl
1 #
2 # 64-bit system call numbers and entry vectors
3 #
4 # The format is:
5 # <number> <abi> <name> <entry point>
6 #
7 # The __x64_sys_*() stubs are created on-the-fly for sys_*() system calls
8 #
9 # The abi is "common", "64" or "x32" for this file.
10 #
11 0      common  read      sys_read
12 1      common  write     sys_write
13 2      common  open      sys_open
```

# sys\_call\_table

The `syscall_64.tbl` will be translated to an array of function pointers (`sys_call_table`) on kernel build

- `scripts/syscalltbl.sh`

```
/* arch/x86/entry/syscall_64.c */  
asmlinkage const sys_call_ptr_t sys_call_table[] =  
{  
#include <asm/syscalls_64.h>  
};
```

File `arch/x86/include/generated/asm/syscalls_64.h` will be generated after kernel build

# Syscall implementation

arch/x86/entry/syscalls/syscall\_64.tbl

```
# 64-bit system call numbers and entry vectors
#
# The format is:
# <number> <abi> <name> <entry point>
#
# The __x64_sys_*() stubs are created on-the-fly for sys_*() system calls
#
# The abi is "common", "64" or "x32" for this file.
#
0      common  read      sys_read
1      common  write     sys_write
2      common  open      sys_open
```

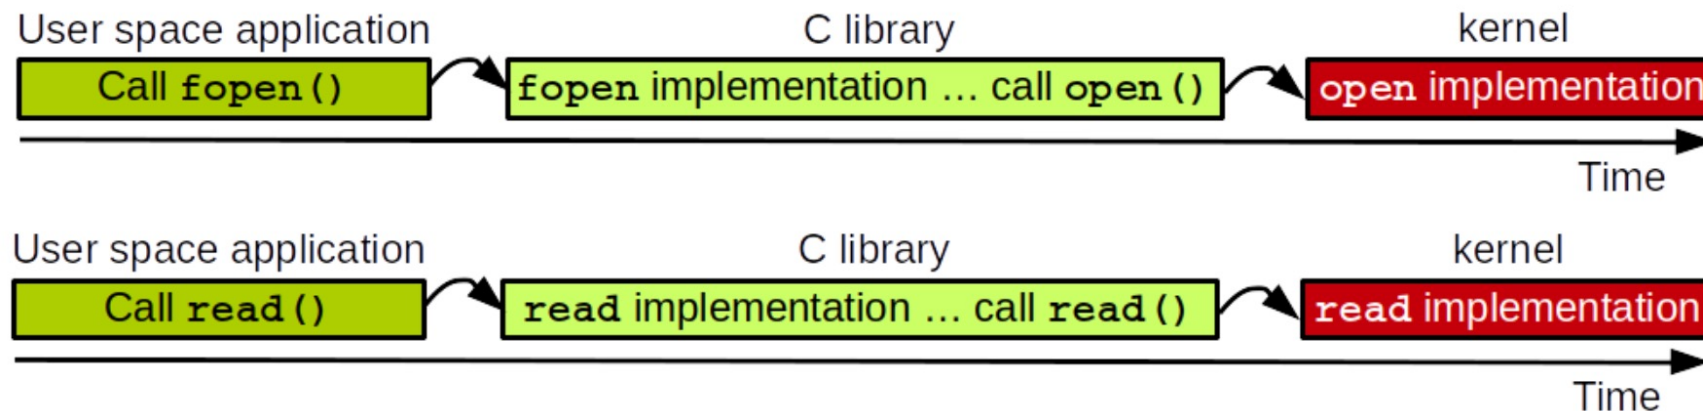
fs/read\_write.c

```
SYSCALL_DEFINE3(read, unsigned int, fd, char __user *, buf, size_t, count)
{
    return ksys_read(fd, buf, count);
}
```

# Invoke a syscall

Syscalls are rarely invoked directly

- Most of them are wrapped by the C library (libc, POSIX API)



# Invoke a syscall

A syscall can be directly called through `syscall()`

- See **man syscall** → A C library function to directly call syscalls

```
#include <unistd.h>
#include <sys/syscall.h> /* for SYS_xxx definitions */
int main(void)
{
    char msg[] = "Hello, world!\n";
    ssize_t bytes_written; /* ssize_t write(int fd, const void *msg, size_t count); */
    bytes_written = syscall(1, 1, msg, 14);
    /*          \      \                                     */
    /*          \      \  +-- fd: standard output           */
    /*          \      \  +-- write syscall id (or SYS_write) */
    return 0;
}
```



# System call instructions

x86 instruction for system call

- `int 0x80`: raise a software interrupt 128 (old)
- `sysenter`: fast system call (x86\_32)
- `syscall`: fast system call (x86\_64)

Passing a syscall ID and parameters

- syscall ID: `%rax`
- parameters (x86\_64): `%rdi`, `%rsi`, `%rdx`, `%r10`, `%r8` and `%r9`
- If a function has more than six arguments, other parameters will be placed on the stack.

# Invoke a syscall

x86\_64 architecture has a `syscall` instruction

```
► cat hello-asm.S  
.data
```

```
msg:  
    .ascii "Hello, world!\n"  
    len = . - msg
```

```
.text  
    .global _start
```

```
_start:  
    mov  $1, %rax    # syscall id: write  
    mov  $1, %rdi    # 1st arg: fd (standard output)  
    mov  $msg, %rsi   # 2nd arg: msg  
    mov  $len, %rdx   # 3rd arg: length of msg  
    syscall          # switch from user space to kernel space  
  
    mov  $60, %rax    # syscall id: exit  
    xor  %rdi, %rdi   # 1st arg: 0  
    syscall          # switch from user space to kernel space
```

# Handling the syscall interrupt

The kernel syscall interrupt handler, system call handler

- `entry_SYSCALL_64` at `arch/x86/entry/entry_64.S`

`entry_SYSCALL_64` is registered at CPU initialization time

- A handler of `syscall` is specified at a `IA32_LSTAR` MSR register
- The address of `IA32_LSTAR` MSR is set to `entry_SYSCALL_64` at boot time: `syscall_init()` at `arch/x86/kernel/cpu/common.c`

# Handling the syscall interrupt

entry\_SYSCALL\_64 invokes the entry function for the syscall ID

- In arch/x86/entry/entry\_64.S
- call do\_syscall\_64
- regs->ax = sys\_call\_table[nr](regs);

```
/* arch/x86/entry/syscall_64.c */
asmlinkage const sys_call_ptr_t sys_call_table[] = {
    [0 ... __NR_syscall_max] = &sys_ni_syscall,
    [0] = sys_read,
    [1] = sys_write,
    ... ..
};
```

# Return from the syscall

---

x86 instruction for system call

- `iret`: interrupt return (x86-32 bit, old)
- `sysexit`: fast return from fast system call (x86-32 bit)
- `sysret`: return from fast system call (x86-64 bit)

# Syscall example -- gettimeofday()

## man gettimeofday

- Get the time

---

GETTIMEOFDAY(2)                      Linux Programmer's Manual                      GETTIMEOFDAY(2)

**NAME**  
gettimeofday, settimeofday - get / set time

**SYNOPSIS**  
`#include <sys/time.h>`

`int gettimeofday(struct timeval *tv, struct timezone *tz);`

`int settimeofday(const struct timeval *tv, const struct timezone *tz);`

Feature Test Macro Requirements for glibc (see `feature_test_macros(7)`):

`settimeofday():`  
Since glibc 2.19:  
    \_DEFAULT\_SOURCE  
Glibc 2.19 and earlier:  
    \_BSD\_SOURCE

**DESCRIPTION**  
The functions `gettimeofday()` and `settimeofday()` can get and set the time as well as a timezone.

The `tv` argument is a `struct timeval` (as specified in `<sys/time.h>`):

```
struct timeval {
    time_t      tv_sec;          /* seconds */
    suseconds_t tv_usec;        /* microseconds */
};
```

# Example C code

```
1 #include <stdio.h>
2 #include <stdlib.h>
3 #include <sys/time.h>
4
5 int main(void)
6 {
7     struct timeval tv;
8     int ret;
9
10    ret = gettimeofday(&tv, NULL);
11    if(ret == -1)
12    {
13        perror("gettimeofday");
14        return EXIT_FAILURE;
15    }
16
17    printf("Local time:\n");
18    printf("  sec:%lu\n", tv.tv_sec);
19    printf("  usec:%lu\n", tv.tv_usec);
20
21    return EXIT_SUCCESS;
22 }
```

time.c

# Kernel implementation

```
SYSCALL_DEFINE2(gettimeofday, struct __kernel_old_timeval __user *, tv,
                 struct timezone __user *, tz)
{
    if (likely(tv != NULL)) {
        struct timespec64 ts;

        ktime_get_real_ts64(&ts);
        if (put_user(ts.tv_sec, &tv->tv_sec) ||
            put_user(ts.tv_nsec / 1000, &tv->tv_usec))
            return -EFAULT;
    }
    if (unlikely(tz != NULL)) {
        if (copy_to_user(tz, &sys_tz, sizeof(sys_tz)))
            return -EFAULT;
    }
    return 0;
}
```

kernel/time/time.c



# User-space vs. kernel-space memory

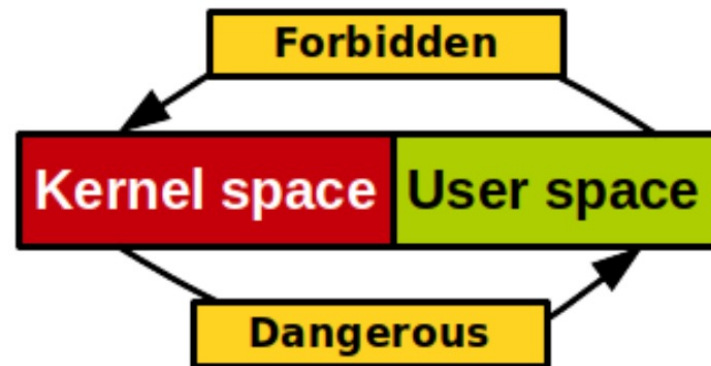
User space cannot access kernel memory

Kernel code must never blindly follow a pointer into user-space

- Q: Why?

Q: How to prevent a user-space access kernel-space memory?

Q: How to safely access user-space memory?



# copy\_{from | to}\_user

*/\* copy user-space memory to kernel-space memory \*/*

**static inline long** copy\_from\_user(**void** \*to, **const void** \_\_user \*from, **unsigned long** n);

*/\* copy kernel-space memory to user-space memory \*/*

**static inline long** copy\_to\_user(**void** \_\_user \*to, **const void** \*from, **unsigned long** n);

Make sure the user-space memory is legitimate

- raise an error if not

... and exist

- wait for user-space memory to swap in

# Implement a new system call

---

1. Write your syscall function
  - Add to the existing file or create a new file
  - Add your new file into the kernel Makefile
2. Add it to the syscall table and assign an ID
  - `arch/x86/entry/syscalls/syscall_64.tbl`
3. Add its prototype in `include/linux/syscalls.h`
4. Compile, reboot, and run
  - Touching the syscall table will trigger the entire kernel compilation

# Implement a new system call

Example: syscall implemented in linux sources in my\_syscall/my\_func.c

Create a linux/my\_syscall/Makefile

- `obj-y += my_func.o`

Add my\_syscall in linux/Makefile

- `core-y += kernel/ certs/ mm/ fs/ ipc/ security/ crypto/  
block/ my_syscall/`

# Why not add a new syscall

**Pros:** Easy to implement and use, fast

**Cons:**

- Needs an official syscall number
- Interface cannot change after implementation
- Must be registered for each architecture
- Probably too much work for small exchanges of information

**Alternative:**

- Create a device node and `read()` and `write()`
- Use `ioctl()`

# Improving syscall performance

System call performance is critical in many applications

- Web server: `select()`, `poll()`
- Game engine: `gettimeofday()`

**Hardware:** add a new fast system call instruction

- `int 0x80` → `syscall`

# Improving syscall performance

**Software:** vDSO (virtual dynamically linked shared object)

- A kernel mechanism for exporting a kernel space routines to user space applications
- No context switching overhead
- E.g., `gettimeofday()`
  - the kernel allows the page containing the current time to be mapped read-only into user space

**Software:** [FlexSC: Exception-less system call, OSDI 2010](#)



# Summary

---

Isolation: CPU privilege on the x86 architecture

System calls: interface for applications to request OS services

Linux system calls: syscall table, syscall handler, and add a syscall

Improve syscall performance



# Next steps

---

Paper reading assignment 1 (due this Friday)

- FlexSC: Exception-less system call, OSDI 2010
  - Summary of the paper; do I like it, or dislike it? Why?
  - Strength of the paper
  - Weakness of the proposed approach
  - Questions/comments if any
- Tips to read a research paper:
  - Watch the presentation video first (if available)
  - Read abstract, introduction and evaluation first; read design/implementation with questions in your mind

# Next steps

---

Hw2 due this Friday

- Linux kernel compilation and boot in a QEMU VM

Hw3 is released (due Jan 26<sup>th</sup>)

- Modify the Linux kernel (~2 hours)
- On top of the QEMU VM from hw2

# Next lecture

---

FlexSC: Exception-less system call

- Optimizing system call performance on multi-core systems
- “We show how FlexSC improves performance of Apache by up to 116%, MySQL by up to 40%, and BIND by up to 105% while requiring no modifications to the applications.”

Kernel data structures

# Further reading

---

LWN: Anatomy of a system call: [part 1](#) and [part 2](#)

LWN: [On vsyscalls and the vDSO](#)

# Feedback

---

